

softmax

January 11, 2016

1 Softmax exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

This exercise is analogous to the SVM exercise. You will:

- implement a fully-vectorized **loss function** for the Softmax classifier
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** with numerical gradient
- use a validation set to **tune the learning rate and regularization strength**
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

```
In [2]: import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading extenrnal modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

```
In [3]: def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the linear classifier. These are the same steps as we used for the
    SVM, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'
    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # subsample the data
    mask = range(num_training, num_training + num_validation)
    X_val = X_train[mask]
```

```

y_val = y_train[mask]
mask = range(num_training)
X_train = X_train[mask]
y_train = y_train[mask]
mask = range(num_test)
X_test = X_test[mask]
y_test = y_test[mask]

# Preprocessing: reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))

# Normalize the data: subtract the mean image
mean_image = np.mean(X_train, axis = 0)
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image

# add bias dimension and transform into columns
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))]).T
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))]).T
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))]).T

return X_train, y_train, X_val, y_val, X_test, y_test

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()
print 'Train data shape: ', X_train.shape
print 'Train labels shape: ', y_train.shape
print 'Validation data shape: ', X_val.shape
print 'Validation labels shape: ', y_val.shape
print 'Test data shape: ', X_test.shape
print 'Test labels shape: ', y_test.shape

```

```

Train data shape: (3073, 49000)
Train labels shape: (49000,)
Validation data shape: (3073, 1000)
Validation labels shape: (1000,)
Test data shape: (3073, 1000)
Test labels shape: (1000,)

```

1.1 Softmax Classifier

Your code for this section will all be written inside **cs231n/classifiers/softmax.py**.

In [4]: *# First implement the naive softmax loss function with nested loops.*
Open the file cs231n/classifiers/softmax.py and implement the
softmax_loss_naive function.

```

from cs231n.classifiers.softmax import softmax_loss_naive
import time

```

Generate a random softmax weight matrix and use it to compute the loss.

```

W = np.random.randn(10, 3073) * 0.0001
loss, grad = softmax_loss_naive(W, X_train, y_train, 0.0)

# As a rough sanity check, our loss should be something close to -log(0.1).
print 'loss: %f' % loss
print 'sanity check: %f' % (-np.log(0.1))

loss: 2.352954
sanity check: 2.302585

```

1.2 Inline Question 1:

Why do we expect our loss to be close to $-\log(0.1)$? Explain briefly.**

Your answer: *Fill this in*

```

In [5]: # Complete the implementation of softmax_loss_naive and implement a (naive)
# version of the gradient that uses nested loops.
loss, grad = softmax_loss_naive(W, X_train, y_train, 0.0)

# As we did for the SVM, use numeric gradient checking as a debugging tool.
# The numeric gradient should be close to the analytic gradient.
from cs231n.gradient_check import grad_check_sparse
f = lambda w: softmax_loss_naive(w, X_train, y_train, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)

numerical: -0.363523 analytic: -0.363523, relative error: 4.031149e-08
numerical: -0.947558 analytic: -0.947558, relative error: 2.887980e-08
numerical: 0.406261 analytic: 0.406261, relative error: 1.465731e-08
numerical: 2.773619 analytic: 2.773619, relative error: 1.267133e-08
numerical: 2.400133 analytic: 2.400133, relative error: 3.195586e-08
numerical: 0.578811 analytic: 0.578811, relative error: 5.938612e-08
numerical: -1.193488 analytic: -1.193488, relative error: 2.080628e-08
numerical: -1.995289 analytic: -1.995289, relative error: 1.341803e-08
numerical: 1.977778 analytic: 1.977778, relative error: 1.310777e-08
numerical: 0.849221 analytic: 0.849221, relative error: 6.014215e-09

In [6]: # Now that we have a naive implementation of the softmax loss function and its gradient,
# implement a vectorized version in softmax_loss_vectorized.
# The two versions should compute the same results, but the vectorized version should be
# much faster.
tic = time.time()
loss_naive, grad_naive = softmax_loss_naive(W, X_train, y_train, 0.00001)
toc = time.time()
print 'naive loss: %e computed in %fs' % (loss_naive, toc - tic)

from cs231n.classifiers.softmax import softmax_loss_vectorized
tic = time.time()
loss_vectorized, grad_vectorized = softmax_loss_vectorized(W, X_train, y_train, 0.00001)
toc = time.time()
print 'vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic)

# As we did for the SVM, we use the Frobenius norm to compare the two versions
# of the gradient.
grad_difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
print 'Loss difference: %f' % np.abs(loss_naive - loss_vectorized)
print 'Gradient difference: %f' % grad_difference

```

naive loss: 2.352954e+00 computed in 0.304337s
vectorized loss: 2.352954e+00 computed in 0.275674s
Loss difference: 0.000000
Gradient difference: 0.000000

```
In [7]: # Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rates and regularization strengths; if you are careful you should be able to
# get a classification accuracy of over 0.35 on the validation set.
from cs231n.classifiers import Softmax
results = {}
best_val = -1
best_softmax = None
learning_rates = [7e-8, 3e-7]
regularization_strengths = [1e4, 3e4]

#####
# TODO:
# Use the validation set to set the learning rate and regularization strength. #
# This should be identical to the validation that you did for the SVM; save   #
# the best trained softmax classifier in best_softmax.                         #
#####
import itertools as it
import copy

from cs231n.classifiers import Softmax

for l_r, reg in it.product(learning_rates, regularization_strengths):
    softmax = Softmax()
    loss_hist = softmax.train(X_train, y_train, learning_rate=l_r, reg=reg,
                              num_iters=1500, verbose=True)
    y_train_pred = softmax.predict(X_train)
    y_val_pred = softmax.predict(X_val)

    train_acc = np.mean(y_train == y_train_pred)
    val_acc = np.mean(y_val == y_val_pred)

    if val_acc > best_val:
        best_val = val_acc
        best_softmax = copy.copy(softmax)

    results[(l_r, reg)] = (train_acc, val_acc)

#####
#                               END OF YOUR CODE                               #
#####

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print 'lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy)

print 'best validation accuracy achieved during cross-validation: %f' % best_val
```

iteration 0 / 1500: loss 160.079160
iteration 100 / 1500: loss 137.928913
iteration 200 / 1500: loss 119.774292
iteration 300 / 1500: loss 104.158282
iteration 400 / 1500: loss 90.664304
iteration 500 / 1500: loss 78.647715
iteration 600 / 1500: loss 68.583049
iteration 700 / 1500: loss 59.801941
iteration 800 / 1500: loss 52.253669
iteration 900 / 1500: loss 45.581820
iteration 1000 / 1500: loss 39.765775
iteration 1100 / 1500: loss 34.840295
iteration 1200 / 1500: loss 30.392724
iteration 1300 / 1500: loss 26.617322
iteration 1400 / 1500: loss 23.483753
iteration 0 / 1500: loss 466.043294
iteration 100 / 1500: loss 305.406777
iteration 200 / 1500: loss 200.728899
iteration 300 / 1500: loss 132.227693
iteration 400 / 1500: loss 87.567721
iteration 500 / 1500: loss 58.078397
iteration 600 / 1500: loss 38.844220
iteration 700 / 1500: loss 26.039527
iteration 800 / 1500: loss 17.856277
iteration 900 / 1500: loss 12.331856
iteration 1000 / 1500: loss 8.808517
iteration 1100 / 1500: loss 6.470089
iteration 1200 / 1500: loss 4.918743
iteration 1300 / 1500: loss 3.945048
iteration 1400 / 1500: loss 3.345241
iteration 0 / 1500: loss 159.190927
iteration 100 / 1500: loss 86.733656
iteration 200 / 1500: loss 48.267235
iteration 300 / 1500: loss 27.033462
iteration 400 / 1500: loss 15.642430
iteration 500 / 1500: loss 9.404781
iteration 600 / 1500: loss 6.070221
iteration 700 / 1500: loss 4.159821
iteration 800 / 1500: loss 3.203543
iteration 900 / 1500: loss 2.601708
iteration 1000 / 1500: loss 2.264816
iteration 1100 / 1500: loss 2.103841
iteration 1200 / 1500: loss 2.033820
iteration 1300 / 1500: loss 2.050385
iteration 1400 / 1500: loss 1.949405
iteration 0 / 1500: loss 464.241736
iteration 100 / 1500: loss 77.085006
iteration 200 / 1500: loss 14.259585
iteration 300 / 1500: loss 4.059058
iteration 400 / 1500: loss 2.340973
iteration 500 / 1500: loss 2.091586
iteration 600 / 1500: loss 2.039393
iteration 700 / 1500: loss 2.089262
iteration 800 / 1500: loss 2.027608

```

iteration 900 / 1500: loss 1.991941
iteration 1000 / 1500: loss 2.062760
iteration 1100 / 1500: loss 2.065556
iteration 1200 / 1500: loss 2.008092
iteration 1300 / 1500: loss 2.034195
iteration 1400 / 1500: loss 1.991874
lr 7.000000e-08 reg 1.000000e+04 train accuracy: 0.293714 val accuracy: 0.298000
lr 7.000000e-08 reg 3.000000e+04 train accuracy: 0.343755 val accuracy: 0.361000
lr 3.000000e-07 reg 1.000000e+04 train accuracy: 0.376490 val accuracy: 0.389000
lr 3.000000e-07 reg 3.000000e+04 train accuracy: 0.340000 val accuracy: 0.352000
best validation accuracy achieved during cross-validation: 0.389000

```

```

In [8]: # evaluate on test set
        # Evaluate the best sum on test set
        y_test_pred = best_softmax.predict(X_test)
        test_accuracy = np.mean(y_test == y_test_pred)
        print 'softmax on raw pixels final test set accuracy: %f' % (test_accuracy, )

```

```

softmax on raw pixels final test set accuracy: 0.361000

```

```


In [9]: # Visualize the learned weights for each class
        w = best_softmax.W[:, :-1] # strip out the bias
        w = w.reshape(10, 32, 32, 3)

        w_min, w_max = np.min(w), np.max(w)

        classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
        for i in xrange(10):
            plt.subplot(2, 5, i + 1)

            # Rescale the weights to be between 0 and 255
            wimg = 255.0 * (w[i].squeeze() - w_min) / (w_max - w_min)
            plt.imshow(wimg.astype('uint8'))
            plt.axis('off')
            plt.title(classes[i])

```



In []: