

convnet

January 12, 2016

1 Train a ConvNet!

We now have a generic solver and a bunch of modularized layers. It's time to put it all together, and train a ConvNet to recognize the classes in CIFAR-10. In this notebook we will walk you through training a simple two-layer ConvNet and then set you free to build the best net that you can to perform well on CIFAR-10.

Open up the file `cs231n/classifiers/convnet.py`; you will see that the `two_layer_convnet` function computes the loss and gradients for a two-layer ConvNet. Note that this function uses the “sandwich” layers defined in `cs231n/layer_utils.py`.

In [6]: *# As usual, a bit of setup*

```
import numpy as np
import matplotlib.pyplot as plt
from cs231n.classifier_trainer import ClassifierTrainer
from cs231n.gradient_check import eval_numerical_gradient
from cs231n.classifiers.convnet import *

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

In [7]: `from cs231n.data_utils import load_CIFAR10`

```
def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the two-layer neural net classifier. These are the same steps as
    we used for the SVM, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'
    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)
```

```

# Subsample the data
mask = range(num_training, num_training + num_validation)
X_val = X_train[mask]
y_val = y_train[mask]
mask = range(num_training)
X_train = X_train[mask]
y_train = y_train[mask]
mask = range(num_test)
X_test = X_test[mask]
y_test = y_test[mask]

# Normalize the data: subtract the mean image
mean_image = np.mean(X_train, axis=0)
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image

# Transpose so that channels come first
X_train = X_train.transpose(0, 3, 1, 2).copy()
X_val = X_val.transpose(0, 3, 1, 2).copy()
x_test = X_test.transpose(0, 3, 1, 2).copy()

return X_train, y_train, X_val, y_val, X_test, y_test

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()
print 'Train data shape: ', X_train.shape
print 'Train labels shape: ', y_train.shape
print 'Validation data shape: ', X_val.shape
print 'Validation labels shape: ', y_val.shape
print 'Test data shape: ', X_test.shape
print 'Test labels shape: ', y_test.shape

```

```

Train data shape: (49000, 3, 32, 32)
Train labels shape: (49000,)
Validation data shape: (1000, 3, 32, 32)
Validation labels shape: (1000,)
Test data shape: (1000, 32, 32, 3)
Test labels shape: (1000,)

```

2 Sanity check loss

After you build a new network, one of the first things you should do is sanity check the loss. When we use the softmax loss, we expect the loss for random weights (and no regularization) to be about $\log(C)$ for C classes. When we add regularization this should go up.

```

In [8]: model = init_two_layer_convnet()

X = np.random.randn(100, 3, 32, 32)
y = np.random.randint(10, size=100)

loss, _ = two_layer_convnet(X, model, y, reg=0)

```

```

# Sanity check: Loss should be about  $\log(10) = 2.3026$ 
print 'Sanity check loss (no regularization): ', loss

# Sanity check: Loss should go up when you add regularization
loss, _ = two_layer_convnet(X, model, y, reg=1)
print 'Sanity check loss (with regularization): ', loss

```

```

Sanity check loss (no regularization): 2.30261884083
Sanity check loss (with regularization): 2.34457328523

```

3 Gradient check

After the loss looks reasonable, you should always use numeric gradient checking to make sure that your backward pass is correct. When you use numeric gradient checking you should use a small amount of artificial data and a small number of neurons at each layer.

```

In [9]: num_inputs = 2
        input_shape = (3, 16, 16)
        reg = 0.0
        num_classes = 10
        X = np.random.randn(num_inputs, *input_shape)
        y = np.random.randint(num_classes, size=num_inputs)

        model = init_two_layer_convnet(num_filters=3, filter_size=3, input_shape=input_shape)
        loss, grads = two_layer_convnet(X, model, y)
        for param_name in sorted(grads):
            f = lambda _: two_layer_convnet(X, model, y)[0]
            param_grad_num = eval_numerical_gradient(f, model[param_name], verbose=False, h=1e-6)
            e = rel_error(param_grad_num, grads[param_name])
            print '%s max relative error: %e' % (param_name, rel_error(param_grad_num, grads[param_name]))

W1 max relative error: 2.695521e-06
W2 max relative error: 1.209239e-05
b1 max relative error: 1.591874e-08
b2 max relative error: 1.399114e-09

```

4 Overfit small data

A nice trick is to train your model with just a few training samples. You should be able to overfit small datasets, which will result in very high training accuracy and comparatively low validation accuracy.

```

In [10]: # Use a two-layer ConvNet to overfit 50 training examples.

        model = init_two_layer_convnet()
        trainer = ClassifierTrainer()
        best_model, loss_history, train_acc_history, val_acc_history = trainer.train(
            X_train[:50], y_train[:50], X_val, y_val, model, two_layer_convnet,
            reg=0.001, momentum=0.9, learning_rate=0.0001, batch_size=10, num_epochs=10,
            verbose=True)

starting iteration 0
Finished epoch 0 / 10: cost 2.299339, train: 0.180000, val 0.076000, lr 1.000000e-04
Finished epoch 1 / 10: cost 2.279487, train: 0.240000, val 0.109000, lr 9.500000e-05
Finished epoch 2 / 10: cost 2.456843, train: 0.280000, val 0.116000, lr 9.025000e-05

```

```

starting iteration 10
Finished epoch 3 / 10: cost 1.863759, train: 0.340000, val 0.163000, lr 8.573750e-05
Finished epoch 4 / 10: cost 1.198169, train: 0.580000, val 0.127000, lr 8.145062e-05
starting iteration 20
Finished epoch 5 / 10: cost 0.999505, train: 0.500000, val 0.186000, lr 7.737809e-05
Finished epoch 6 / 10: cost 0.829464, train: 0.600000, val 0.185000, lr 7.350919e-05
starting iteration 30
Finished epoch 7 / 10: cost 0.713109, train: 0.740000, val 0.158000, lr 6.983373e-05
Finished epoch 8 / 10: cost 0.758027, train: 0.820000, val 0.136000, lr 6.634204e-05
starting iteration 40
Finished epoch 9 / 10: cost 0.395614, train: 0.880000, val 0.189000, lr 6.302494e-05
Finished epoch 10 / 10: cost 0.185496, train: 0.940000, val 0.180000, lr 5.987369e-05
finished optimization. best validation accuracy: 0.189000

```


Plotting the loss, training accuracy, and validation accuracy should show clear overfitting:

```

In [11]: plt.subplot(2, 1, 1)
         plt.plot(loss_history)
         plt.xlabel('iteration')
         plt.ylabel('loss')

         plt.subplot(2, 1, 2)
         plt.plot(train_acc_history)
         plt.plot(val_acc_history)
         plt.legend(['train', 'val'], loc='upper left')
         plt.xlabel('epoch')
         plt.ylabel('accuracy')
         plt.show()

```



5 Train the net

Once the above works, training the net is the next thing to try. You can set the `acc_frequency` parameter to change the frequency at which the training and validation set accuracies are tested. If your parameters are set properly, you should see the training and validation accuracy start to improve within a hundred iterations, and you should be able to train a reasonable model with just one epoch.

Using the parameters below you should be able to get around 50% accuracy on the validation set.

```
In [12]: model = init_two_layer_convnet(filter_size=7)
         trainer = ClassifierTrainer()
         best_model, loss_history, train_acc_history, val_acc_history = trainer.train(
             X_train, y_train, X_val, y_val, model, two_layer_convnet,
             reg=0.001, momentum=0.9, learning_rate=0.0001, batch_size=50, num_epochs=1,
             acc_frequency=50, verbose=True)

starting iteration  0
Finished epoch 0 / 1: cost 2.298741, train: 0.106000, val 0.108000, lr 1.000000e-04
starting iteration 10
starting iteration 20
starting iteration 30
starting iteration 40
starting iteration 50
```

Finished epoch 0 / 1: cost 1.934964, train: 0.319000, val 0.336000, lr 1.000000e-04
starting iteration 60
starting iteration 70
starting iteration 80
starting iteration 90
starting iteration 100
Finished epoch 0 / 1: cost 1.947233, train: 0.289000, val 0.308000, lr 1.000000e-04
starting iteration 110
starting iteration 120
starting iteration 130
starting iteration 140
starting iteration 150
Finished epoch 0 / 1: cost 1.616109, train: 0.403000, val 0.379000, lr 1.000000e-04
starting iteration 160
starting iteration 170
starting iteration 180
starting iteration 190
starting iteration 200
Finished epoch 0 / 1: cost 1.689934, train: 0.420000, val 0.415000, lr 1.000000e-04
starting iteration 210
starting iteration 220
starting iteration 230
starting iteration 240
starting iteration 250
Finished epoch 0 / 1: cost 1.461824, train: 0.414000, val 0.408000, lr 1.000000e-04
starting iteration 260
starting iteration 270
starting iteration 280
starting iteration 290
starting iteration 300
Finished epoch 0 / 1: cost 1.394433, train: 0.465000, val 0.450000, lr 1.000000e-04
starting iteration 310
starting iteration 320
starting iteration 330
starting iteration 340
starting iteration 350
Finished epoch 0 / 1: cost 1.366479, train: 0.426000, val 0.410000, lr 1.000000e-04
starting iteration 360
starting iteration 370
starting iteration 380
starting iteration 390
starting iteration 400
Finished epoch 0 / 1: cost 1.698907, train: 0.408000, val 0.432000, lr 1.000000e-04
starting iteration 410
starting iteration 420
starting iteration 430
starting iteration 440
starting iteration 450
Finished epoch 0 / 1: cost 1.808498, train: 0.448000, val 0.413000, lr 1.000000e-04
starting iteration 460
starting iteration 470
starting iteration 480
starting iteration 490
starting iteration 500

Finished epoch 0 / 1: cost 2.835956, train: 0.472000, val 0.440000, lr 1.000000e-04
starting iteration 510
starting iteration 520
starting iteration 530
starting iteration 540
starting iteration 550
Finished epoch 0 / 1: cost 1.945160, train: 0.443000, val 0.458000, lr 1.000000e-04
starting iteration 560
starting iteration 570
starting iteration 580
starting iteration 590
starting iteration 600
Finished epoch 0 / 1: cost 1.716514, train: 0.499000, val 0.475000, lr 1.000000e-04
starting iteration 610
starting iteration 620
starting iteration 630
starting iteration 640
starting iteration 650
Finished epoch 0 / 1: cost 1.739311, train: 0.515000, val 0.471000, lr 1.000000e-04
starting iteration 660
starting iteration 670
starting iteration 680
starting iteration 690
starting iteration 700
Finished epoch 0 / 1: cost 1.226189, train: 0.482000, val 0.479000, lr 1.000000e-04
starting iteration 710
starting iteration 720
starting iteration 730
starting iteration 740
starting iteration 750
Finished epoch 0 / 1: cost 1.617910, train: 0.535000, val 0.494000, lr 1.000000e-04
starting iteration 760
starting iteration 770
starting iteration 780
starting iteration 790
starting iteration 800
Finished epoch 0 / 1: cost 1.438191, train: 0.475000, val 0.457000, lr 1.000000e-04
starting iteration 810
starting iteration 820
starting iteration 830
starting iteration 840
starting iteration 850
Finished epoch 0 / 1: cost 2.431353, train: 0.438000, val 0.442000, lr 1.000000e-04
starting iteration 860
starting iteration 870
starting iteration 880
starting iteration 890
starting iteration 900
Finished epoch 0 / 1: cost 2.309801, train: 0.528000, val 0.519000, lr 1.000000e-04
starting iteration 910
starting iteration 920
starting iteration 930
starting iteration 940
starting iteration 950

```
Finished epoch 0 / 1: cost 1.674035, train: 0.473000, val 0.443000, lr 1.000000e-04
starting iteration 960
starting iteration 970
Finished epoch 1 / 1: cost 2.045021, train: 0.491000, val 0.463000, lr 9.500000e-05
finished optimization. best validation accuracy: 0.519000
```


6 Visualize weights

We can visualize the convolutional weights from the first layer. If everything worked properly, these will usually be edges and blobs of various colors and orientations.

```
In [8]: from cs231n.vis_utils import visualize_grid
```

```
grid = visualize_grid(best_model['W1'].transpose(0, 2, 3, 1))
plt.imshow(grid.astype('uint8'))
```

```
Out[8]: <matplotlib.image.AxesImage at 0x1069d21d0>
```



7 Experiment!

Experiment and try to get the best performance that you can on CIFAR-10 using a ConvNet. Here are some ideas to get you started:

7.0.1 Things you should try:

- Filter size: Above we used 7x7; this makes pretty pictures but smaller filters may be more efficient
- Number of filters: Above we used 32 filters. Do more or fewer do better?
- Network depth: The network above has two layers of trainable parameters. Can you do better with a deeper network? You can implement alternative architectures in the file `cs231n/classifiers/convnet.py`. Some good architectures to try include:
 - [conv-relu-pool]xN - conv - relu - [affine]xM - [softmax or SVM]
 - [conv-relu-pool]XN - [affine]XM - [softmax or SVM]
 - [conv-relu-conv-relu-pool]xN - [affine]xM - [softmax or SVM]

7.0.2 Tips for training

For each network architecture that you try, you should tune the learning rate and regularization strength. When doing this there are a couple important things to keep in mind:

- If the parameters are working well, you should see improvement within a few hundred iterations
- Remember the course-to-fine approach for hyperparameter tuning: start by testing a large range of hyperparameters for just a few training iterations to find the combinations of parameters that are working at all.
- Once you have found some sets of parameters that seem to work, search more finely around these parameters. You may need to train for more epochs.

7.0.3 Going above and beyond

If you are feeling adventurous there are many other features you can implement to try and improve your performance. You are **not required** to implement any of these; however they would be good things to try for extra credit.

- Alternative update steps: For the assignment we implemented SGD+momentum and RMSprop; you could try alternatives like AdaGrad or AdaDelta.
- Other forms of regularization such as L1 or Dropout
- Alternative activation functions such as leaky ReLU or maxout
- Model ensembles
- Data augmentation

7.0.4 What we expect

At the very least, you should be able to train a ConvNet that gets at least 65% accuracy on the validation set. This is just a lower bound - if you are careful it should be possible to get accuracies much higher than that! Extra credit points will be awarded for particularly high-scoring models or unique approaches.

You should use the space below to experiment and train your network. The final cell in this notebook should contain the training, validation, and test set accuracies for your final trained network. In this notebook you should also write an explanation of what you did, any additional features that you implemented, and any visualizations or graphs that you make in the process of training and evaluating your network.

Have fun and happy training!

```

In [35]: # TODO: Train a ConvNet to do really well on CIFAR-10!
         from cs231n.classifiers.convnet import *
         from cs231n.layers import *

         #model = init_two_layer_convnet(filter_size=3, num_filters=64)
         #trainer = ClassifierTrainer()
         #best_model, loss_history, train_acc_history, val_acc_history = trainer.train(
         #      X_train, y_train, X_val, y_val, model, two_layer_convnet,
         #      reg=0.001, momentum=0.9, learning_rate=0.0001, batch_size=50, num_epochs=1,
         #      acc_frequency=50, verbose=True)

         model = init_my_convnet()
         trainer = ClassifierTrainer()
         best_model, loss_history, train_acc_history, val_acc_history = trainer.train(
             X_train, y_train, X_val, y_val, model, my_convnet, momentum=0.9, learning_rate=0.0005,
             batch_size=50, num_epochs=5, acc_frequency=50, verbose=True)

starting iteration  0
Finished epoch 0 / 5: cost 2.300679, train: 0.081000, val 0.102000, lr 5.000000e-04
starting iteration  10
starting iteration  20
starting iteration  30
starting iteration  40
starting iteration  50
Finished epoch 0 / 5: cost 2.200456, train: 0.233000, val 0.251000, lr 5.000000e-04
starting iteration  60
starting iteration  70
starting iteration  80
starting iteration  90
starting iteration 100
Finished epoch 0 / 5: cost 2.011854, train: 0.297000, val 0.319000, lr 5.000000e-04
starting iteration 110
starting iteration 120
starting iteration 130
starting iteration 140
starting iteration 150
Finished epoch 0 / 5: cost 1.868806, train: 0.367000, val 0.375000, lr 5.000000e-04
starting iteration 160
starting iteration 170
starting iteration 180
starting iteration 190
starting iteration 200
Finished epoch 0 / 5: cost 1.982107, train: 0.400000, val 0.404000, lr 5.000000e-04
starting iteration 210
starting iteration 220
starting iteration 230
starting iteration 240
starting iteration 250
Finished epoch 0 / 5: cost 1.626580, train: 0.431000, val 0.444000, lr 5.000000e-04
starting iteration 260
starting iteration 270
starting iteration 280
starting iteration 290
starting iteration 300

```

Finished epoch 0 / 5: cost 1.516094, train: 0.421000, val 0.447000, lr 5.000000e-04
starting iteration 310
starting iteration 320
starting iteration 330
starting iteration 340
starting iteration 350
Finished epoch 0 / 5: cost 1.481755, train: 0.457000, val 0.470000, lr 5.000000e-04
starting iteration 360
starting iteration 370
starting iteration 380
starting iteration 390
starting iteration 400
Finished epoch 0 / 5: cost 1.550152, train: 0.476000, val 0.477000, lr 5.000000e-04
starting iteration 410
starting iteration 420
starting iteration 430
starting iteration 440
starting iteration 450
Finished epoch 0 / 5: cost 1.520160, train: 0.500000, val 0.485000, lr 5.000000e-04
starting iteration 460
starting iteration 470
starting iteration 480
starting iteration 490
starting iteration 500
Finished epoch 0 / 5: cost 1.626071, train: 0.526000, val 0.518000, lr 5.000000e-04
starting iteration 510
starting iteration 520
starting iteration 530
starting iteration 540
starting iteration 550
Finished epoch 0 / 5: cost 1.209234, train: 0.522000, val 0.543000, lr 5.000000e-04
starting iteration 560
starting iteration 570
starting iteration 580
starting iteration 590
starting iteration 600
Finished epoch 0 / 5: cost 1.277498, train: 0.519000, val 0.524000, lr 5.000000e-04
starting iteration 610
starting iteration 620
starting iteration 630
starting iteration 640
starting iteration 650
Finished epoch 0 / 5: cost 1.236845, train: 0.537000, val 0.542000, lr 5.000000e-04
starting iteration 660
starting iteration 670
starting iteration 680
starting iteration 690
starting iteration 700
Finished epoch 0 / 5: cost 1.311758, train: 0.561000, val 0.578000, lr 5.000000e-04
starting iteration 710
starting iteration 720
starting iteration 730
starting iteration 740
starting iteration 750

Finished epoch 0 / 5: cost 1.264987, train: 0.606000, val 0.575000, lr 5.000000e-04
starting iteration 760
starting iteration 770
starting iteration 780
starting iteration 790
starting iteration 800
Finished epoch 0 / 5: cost 1.207748, train: 0.574000, val 0.563000, lr 5.000000e-04
starting iteration 810
starting iteration 820
starting iteration 830
starting iteration 840
starting iteration 850
Finished epoch 0 / 5: cost 1.367678, train: 0.609000, val 0.585000, lr 5.000000e-04
starting iteration 860
starting iteration 870
starting iteration 880
starting iteration 890
starting iteration 900
Finished epoch 0 / 5: cost 1.154976, train: 0.632000, val 0.601000, lr 5.000000e-04
starting iteration 910
starting iteration 920
starting iteration 930
starting iteration 940
starting iteration 950
Finished epoch 0 / 5: cost 1.243287, train: 0.610000, val 0.589000, lr 5.000000e-04
starting iteration 960
starting iteration 970
Finished epoch 1 / 5: cost 1.069391, train: 0.606000, val 0.583000, lr 4.750000e-04
starting iteration 980
starting iteration 990
starting iteration 1000
Finished epoch 1 / 5: cost 1.215283, train: 0.601000, val 0.606000, lr 4.750000e-04
starting iteration 1010
starting iteration 1020
starting iteration 1030
starting iteration 1040
starting iteration 1050
Finished epoch 1 / 5: cost 1.183057, train: 0.592000, val 0.575000, lr 4.750000e-04
starting iteration 1060
starting iteration 1070
starting iteration 1080
starting iteration 1090
starting iteration 1100
Finished epoch 1 / 5: cost 1.375366, train: 0.589000, val 0.584000, lr 4.750000e-04
starting iteration 1110
starting iteration 1120
starting iteration 1130
starting iteration 1140
starting iteration 1150
Finished epoch 1 / 5: cost 1.222407, train: 0.614000, val 0.624000, lr 4.750000e-04
starting iteration 1160
starting iteration 1170
starting iteration 1180
starting iteration 1190

starting iteration 1200
Finished epoch 1 / 5: cost 0.874685, train: 0.638000, val 0.621000, lr 4.750000e-04
starting iteration 1210
starting iteration 1220
starting iteration 1230
starting iteration 1240
starting iteration 1250
Finished epoch 1 / 5: cost 0.868269, train: 0.628000, val 0.607000, lr 4.750000e-04
starting iteration 1260
starting iteration 1270
starting iteration 1280
starting iteration 1290
starting iteration 1300
Finished epoch 1 / 5: cost 0.899256, train: 0.648000, val 0.618000, lr 4.750000e-04
starting iteration 1310
starting iteration 1320
starting iteration 1330
starting iteration 1340
starting iteration 1350
Finished epoch 1 / 5: cost 0.999799, train: 0.632000, val 0.615000, lr 4.750000e-04
starting iteration 1360
starting iteration 1370
starting iteration 1380
starting iteration 1390
starting iteration 1400
Finished epoch 1 / 5: cost 0.885498, train: 0.672000, val 0.636000, lr 4.750000e-04
starting iteration 1410
starting iteration 1420
starting iteration 1430
starting iteration 1440
starting iteration 1450
Finished epoch 1 / 5: cost 1.156575, train: 0.622000, val 0.630000, lr 4.750000e-04
starting iteration 1460
starting iteration 1470
starting iteration 1480
starting iteration 1490
starting iteration 1500
Finished epoch 1 / 5: cost 0.927210, train: 0.651000, val 0.649000, lr 4.750000e-04
starting iteration 1510
starting iteration 1520
starting iteration 1530
starting iteration 1540
starting iteration 1550
Finished epoch 1 / 5: cost 0.686877, train: 0.689000, val 0.651000, lr 4.750000e-04
starting iteration 1560
starting iteration 1570
starting iteration 1580
starting iteration 1590
starting iteration 1600
Finished epoch 1 / 5: cost 0.830908, train: 0.668000, val 0.665000, lr 4.750000e-04
starting iteration 1610
starting iteration 1620
starting iteration 1630
starting iteration 1640

starting iteration 1650
Finished epoch 1 / 5: cost 0.804865, train: 0.691000, val 0.647000, lr 4.750000e-04
starting iteration 1660
starting iteration 1670
starting iteration 1680
starting iteration 1690
starting iteration 1700
Finished epoch 1 / 5: cost 0.969134, train: 0.622000, val 0.654000, lr 4.750000e-04
starting iteration 1710
starting iteration 1720
starting iteration 1730
starting iteration 1740
starting iteration 1750
Finished epoch 1 / 5: cost 0.833219, train: 0.675000, val 0.655000, lr 4.750000e-04
starting iteration 1760
starting iteration 1770
starting iteration 1780
starting iteration 1790
starting iteration 1800
Finished epoch 1 / 5: cost 0.699026, train: 0.636000, val 0.649000, lr 4.750000e-04
starting iteration 1810
starting iteration 1820
starting iteration 1830
starting iteration 1840
starting iteration 1850
Finished epoch 1 / 5: cost 0.931040, train: 0.675000, val 0.632000, lr 4.750000e-04
starting iteration 1860
starting iteration 1870
starting iteration 1880
starting iteration 1890
starting iteration 1900
Finished epoch 1 / 5: cost 0.890078, train: 0.709000, val 0.660000, lr 4.750000e-04
starting iteration 1910
starting iteration 1920
starting iteration 1930
starting iteration 1940
starting iteration 1950
Finished epoch 1 / 5: cost 0.938553, train: 0.693000, val 0.661000, lr 4.750000e-04
Finished epoch 2 / 5: cost 0.901983, train: 0.693000, val 0.658000, lr 4.512500e-04
starting iteration 1960
starting iteration 1970
starting iteration 1980
starting iteration 1990
starting iteration 2000
Finished epoch 2 / 5: cost 0.901837, train: 0.703000, val 0.680000, lr 4.512500e-04
starting iteration 2010
starting iteration 2020
starting iteration 2030
starting iteration 2040
starting iteration 2050
Finished epoch 2 / 5: cost 1.011459, train: 0.725000, val 0.667000, lr 4.512500e-04
starting iteration 2060
starting iteration 2070
starting iteration 2080

starting iteration 2090
starting iteration 2100
Finished epoch 2 / 5: cost 1.004735, train: 0.731000, val 0.683000, lr 4.512500e-04
starting iteration 2110
starting iteration 2120
starting iteration 2130
starting iteration 2140
starting iteration 2150
Finished epoch 2 / 5: cost 1.187211, train: 0.690000, val 0.657000, lr 4.512500e-04
starting iteration 2160
starting iteration 2170
starting iteration 2180
starting iteration 2190
starting iteration 2200
Finished epoch 2 / 5: cost 1.319832, train: 0.715000, val 0.668000, lr 4.512500e-04
starting iteration 2210
starting iteration 2220
starting iteration 2230
starting iteration 2240
starting iteration 2250
Finished epoch 2 / 5: cost 0.560827, train: 0.713000, val 0.668000, lr 4.512500e-04
starting iteration 2260
starting iteration 2270
starting iteration 2280
starting iteration 2290
starting iteration 2300
Finished epoch 2 / 5: cost 0.899339, train: 0.718000, val 0.690000, lr 4.512500e-04
starting iteration 2310
starting iteration 2320
starting iteration 2330
starting iteration 2340
starting iteration 2350
Finished epoch 2 / 5: cost 1.006655, train: 0.713000, val 0.688000, lr 4.512500e-04
starting iteration 2360
starting iteration 2370
starting iteration 2380
starting iteration 2390
starting iteration 2400
Finished epoch 2 / 5: cost 1.024293, train: 0.729000, val 0.683000, lr 4.512500e-04
starting iteration 2410
starting iteration 2420
starting iteration 2430
starting iteration 2440
starting iteration 2450
Finished epoch 2 / 5: cost 0.751033, train: 0.712000, val 0.689000, lr 4.512500e-04
starting iteration 2460
starting iteration 2470
starting iteration 2480
starting iteration 2490
starting iteration 2500
Finished epoch 2 / 5: cost 0.714302, train: 0.708000, val 0.665000, lr 4.512500e-04
starting iteration 2510
starting iteration 2520
starting iteration 2530

starting iteration 2540
starting iteration 2550
Finished epoch 2 / 5: cost 0.957750, train: 0.718000, val 0.693000, lr 4.512500e-04
starting iteration 2560
starting iteration 2570
starting iteration 2580
starting iteration 2590
starting iteration 2600
Finished epoch 2 / 5: cost 0.602884, train: 0.762000, val 0.700000, lr 4.512500e-04
starting iteration 2610
starting iteration 2620
starting iteration 2630
starting iteration 2640
starting iteration 2650
Finished epoch 2 / 5: cost 0.655529, train: 0.736000, val 0.697000, lr 4.512500e-04
starting iteration 2660
starting iteration 2670
starting iteration 2680
starting iteration 2690
starting iteration 2700
Finished epoch 2 / 5: cost 0.707096, train: 0.758000, val 0.710000, lr 4.512500e-04
starting iteration 2710
starting iteration 2720
starting iteration 2730
starting iteration 2740
starting iteration 2750
Finished epoch 2 / 5: cost 0.550205, train: 0.734000, val 0.685000, lr 4.512500e-04
starting iteration 2760
starting iteration 2770
starting iteration 2780
starting iteration 2790
starting iteration 2800
Finished epoch 2 / 5: cost 0.629230, train: 0.755000, val 0.708000, lr 4.512500e-04
starting iteration 2810
starting iteration 2820
starting iteration 2830
starting iteration 2840
starting iteration 2850
Finished epoch 2 / 5: cost 0.965902, train: 0.724000, val 0.677000, lr 4.512500e-04
starting iteration 2860
starting iteration 2870
starting iteration 2880
starting iteration 2890
starting iteration 2900
Finished epoch 2 / 5: cost 0.852758, train: 0.747000, val 0.708000, lr 4.512500e-04
starting iteration 2910
starting iteration 2920
starting iteration 2930
Finished epoch 3 / 5: cost 0.554288, train: 0.735000, val 0.702000, lr 4.286875e-04
starting iteration 2940
starting iteration 2950
Finished epoch 3 / 5: cost 0.606847, train: 0.760000, val 0.699000, lr 4.286875e-04
starting iteration 2960
starting iteration 2970

starting iteration 2980
starting iteration 2990
starting iteration 3000
Finished epoch 3 / 5: cost 0.717710, train: 0.759000, val 0.708000, lr 4.286875e-04
starting iteration 3010
starting iteration 3020
starting iteration 3030
starting iteration 3040
starting iteration 3050
Finished epoch 3 / 5: cost 0.787513, train: 0.763000, val 0.695000, lr 4.286875e-04
starting iteration 3060
starting iteration 3070
starting iteration 3080
starting iteration 3090
starting iteration 3100
Finished epoch 3 / 5: cost 0.704798, train: 0.710000, val 0.667000, lr 4.286875e-04
starting iteration 3110
starting iteration 3120
starting iteration 3130
starting iteration 3140
starting iteration 3150
Finished epoch 3 / 5: cost 0.599914, train: 0.773000, val 0.708000, lr 4.286875e-04
starting iteration 3160
starting iteration 3170
starting iteration 3180
starting iteration 3190
starting iteration 3200
Finished epoch 3 / 5: cost 0.781879, train: 0.764000, val 0.698000, lr 4.286875e-04
starting iteration 3210
starting iteration 3220
starting iteration 3230
starting iteration 3240
starting iteration 3250
Finished epoch 3 / 5: cost 0.639903, train: 0.774000, val 0.713000, lr 4.286875e-04
starting iteration 3260
starting iteration 3270
starting iteration 3280
starting iteration 3290
starting iteration 3300
Finished epoch 3 / 5: cost 0.692868, train: 0.747000, val 0.705000, lr 4.286875e-04
starting iteration 3310
starting iteration 3320
starting iteration 3330
starting iteration 3340
starting iteration 3350
Finished epoch 3 / 5: cost 0.809297, train: 0.762000, val 0.724000, lr 4.286875e-04
starting iteration 3360
starting iteration 3370
starting iteration 3380
starting iteration 3390
starting iteration 3400
Finished epoch 3 / 5: cost 0.844133, train: 0.771000, val 0.720000, lr 4.286875e-04
starting iteration 3410
starting iteration 3420

starting iteration 3430
starting iteration 3440
starting iteration 3450
Finished epoch 3 / 5: cost 0.456535, train: 0.753000, val 0.716000, lr 4.286875e-04
starting iteration 3460
starting iteration 3470
starting iteration 3480
starting iteration 3490
starting iteration 3500
Finished epoch 3 / 5: cost 0.523346, train: 0.776000, val 0.691000, lr 4.286875e-04
starting iteration 3510
starting iteration 3520
starting iteration 3530
starting iteration 3540
starting iteration 3550
Finished epoch 3 / 5: cost 0.522825, train: 0.764000, val 0.706000, lr 4.286875e-04
starting iteration 3560
starting iteration 3570
starting iteration 3580
starting iteration 3590
starting iteration 3600
Finished epoch 3 / 5: cost 0.764020, train: 0.777000, val 0.706000, lr 4.286875e-04
starting iteration 3610
starting iteration 3620
starting iteration 3630
starting iteration 3640
starting iteration 3650
Finished epoch 3 / 5: cost 0.751349, train: 0.799000, val 0.708000, lr 4.286875e-04
starting iteration 3660
starting iteration 3670
starting iteration 3680
starting iteration 3690
starting iteration 3700
Finished epoch 3 / 5: cost 0.715394, train: 0.775000, val 0.704000, lr 4.286875e-04
starting iteration 3710
starting iteration 3720
starting iteration 3730
starting iteration 3740
starting iteration 3750
Finished epoch 3 / 5: cost 0.582312, train: 0.766000, val 0.705000, lr 4.286875e-04
starting iteration 3760
starting iteration 3770
starting iteration 3780
starting iteration 3790
starting iteration 3800
Finished epoch 3 / 5: cost 0.736625, train: 0.811000, val 0.724000, lr 4.286875e-04
starting iteration 3810
starting iteration 3820
starting iteration 3830
starting iteration 3840
starting iteration 3850
Finished epoch 3 / 5: cost 0.599804, train: 0.784000, val 0.709000, lr 4.286875e-04
starting iteration 3860
starting iteration 3870

starting iteration 3880
starting iteration 3890
starting iteration 3900
Finished epoch 3 / 5: cost 0.505844, train: 0.771000, val 0.702000, lr 4.286875e-04
starting iteration 3910
Finished epoch 4 / 5: cost 0.667030, train: 0.797000, val 0.716000, lr 4.072531e-04
starting iteration 3920
starting iteration 3930
starting iteration 3940
starting iteration 3950
Finished epoch 4 / 5: cost 0.430920, train: 0.809000, val 0.711000, lr 4.072531e-04
starting iteration 3960
starting iteration 3970
starting iteration 3980
starting iteration 3990
starting iteration 4000
Finished epoch 4 / 5: cost 0.543429, train: 0.789000, val 0.718000, lr 4.072531e-04
starting iteration 4010
starting iteration 4020
starting iteration 4030
starting iteration 4040
starting iteration 4050
Finished epoch 4 / 5: cost 0.567855, train: 0.801000, val 0.730000, lr 4.072531e-04
starting iteration 4060
starting iteration 4070
starting iteration 4080
starting iteration 4090
starting iteration 4100
Finished epoch 4 / 5: cost 0.563883, train: 0.804000, val 0.722000, lr 4.072531e-04
starting iteration 4110
starting iteration 4120
starting iteration 4130
starting iteration 4140
starting iteration 4150
Finished epoch 4 / 5: cost 0.681282, train: 0.765000, val 0.702000, lr 4.072531e-04
starting iteration 4160
starting iteration 4170
starting iteration 4180
starting iteration 4190
starting iteration 4200
Finished epoch 4 / 5: cost 0.908928, train: 0.792000, val 0.733000, lr 4.072531e-04
starting iteration 4210
starting iteration 4220
starting iteration 4230
starting iteration 4240
starting iteration 4250
Finished epoch 4 / 5: cost 0.891740, train: 0.791000, val 0.732000, lr 4.072531e-04
starting iteration 4260
starting iteration 4270
starting iteration 4280
starting iteration 4290
starting iteration 4300
Finished epoch 4 / 5: cost 0.569775, train: 0.834000, val 0.735000, lr 4.072531e-04
starting iteration 4310

starting iteration 4320
starting iteration 4330
starting iteration 4340
starting iteration 4350
Finished epoch 4 / 5: cost 0.409690, train: 0.802000, val 0.748000, lr 4.072531e-04
starting iteration 4360
starting iteration 4370
starting iteration 4380
starting iteration 4390
starting iteration 4400
Finished epoch 4 / 5: cost 0.634047, train: 0.818000, val 0.715000, lr 4.072531e-04
starting iteration 4410
starting iteration 4420
starting iteration 4430
starting iteration 4440
starting iteration 4450
Finished epoch 4 / 5: cost 0.577226, train: 0.838000, val 0.736000, lr 4.072531e-04
starting iteration 4460
starting iteration 4470
starting iteration 4480
starting iteration 4490
starting iteration 4500
Finished epoch 4 / 5: cost 0.647941, train: 0.792000, val 0.735000, lr 4.072531e-04
starting iteration 4510
starting iteration 4520
starting iteration 4530
starting iteration 4540
starting iteration 4550
Finished epoch 4 / 5: cost 0.578878, train: 0.822000, val 0.736000, lr 4.072531e-04
starting iteration 4560
starting iteration 4570
starting iteration 4580
starting iteration 4590
starting iteration 4600
Finished epoch 4 / 5: cost 0.435067, train: 0.812000, val 0.731000, lr 4.072531e-04
starting iteration 4610
starting iteration 4620
starting iteration 4630
starting iteration 4640
starting iteration 4650
Finished epoch 4 / 5: cost 0.423294, train: 0.811000, val 0.714000, lr 4.072531e-04
starting iteration 4660
starting iteration 4670
starting iteration 4680
starting iteration 4690
starting iteration 4700
Finished epoch 4 / 5: cost 0.714353, train: 0.823000, val 0.734000, lr 4.072531e-04
starting iteration 4710
starting iteration 4720
starting iteration 4730
starting iteration 4740
starting iteration 4750
Finished epoch 4 / 5: cost 0.534575, train: 0.794000, val 0.734000, lr 4.072531e-04
starting iteration 4760

```

starting iteration 4770
starting iteration 4780
starting iteration 4790
starting iteration 4800
Finished epoch 4 / 5: cost 0.375923, train: 0.823000, val 0.745000, lr 4.072531e-04
starting iteration 4810
starting iteration 4820
starting iteration 4830
starting iteration 4840
starting iteration 4850
Finished epoch 4 / 5: cost 0.363778, train: 0.806000, val 0.745000, lr 4.072531e-04
starting iteration 4860
starting iteration 4870
starting iteration 4880
starting iteration 4890
Finished epoch 5 / 5: cost 0.409641, train: 0.823000, val 0.730000, lr 3.868905e-04
finished optimization. best validation accuracy: 0.748000

```

```

In [41]: #NOTE: There is a typo in the get_CIFAR10_data code above
         #right before the return statement.
         #Rather than regenerate this book, I am running this line below
         x_test = X_test.transpose(0, 3, 1, 2).copy()

```

```

In [42]: print X_val.shape, x_test.shape

```

```

(1000, 3, 32, 32) (1000, 3, 32, 32)

```

```

In [43]: #Make sure that I get 0.748 for the validation set with best_model
         scores = my_convnet(X_val, best_model)
         np.mean(np.argmax(scores, axis=1) == y_val)

```

```

Out[43]: 0.748

```

```

In [44]: #Now check the test set score
         scores = my_convnet(x_test, best_model)
         np.mean(np.argmax(scores, axis=1) == y_test)

```

```

Out[44]: 0.72799999999999998

```

```

In [ ]:

```